



Do the models recognize the truth of the items? (10 scenarios × 3 frontier models; R1 n≈20/bar until full re-run)

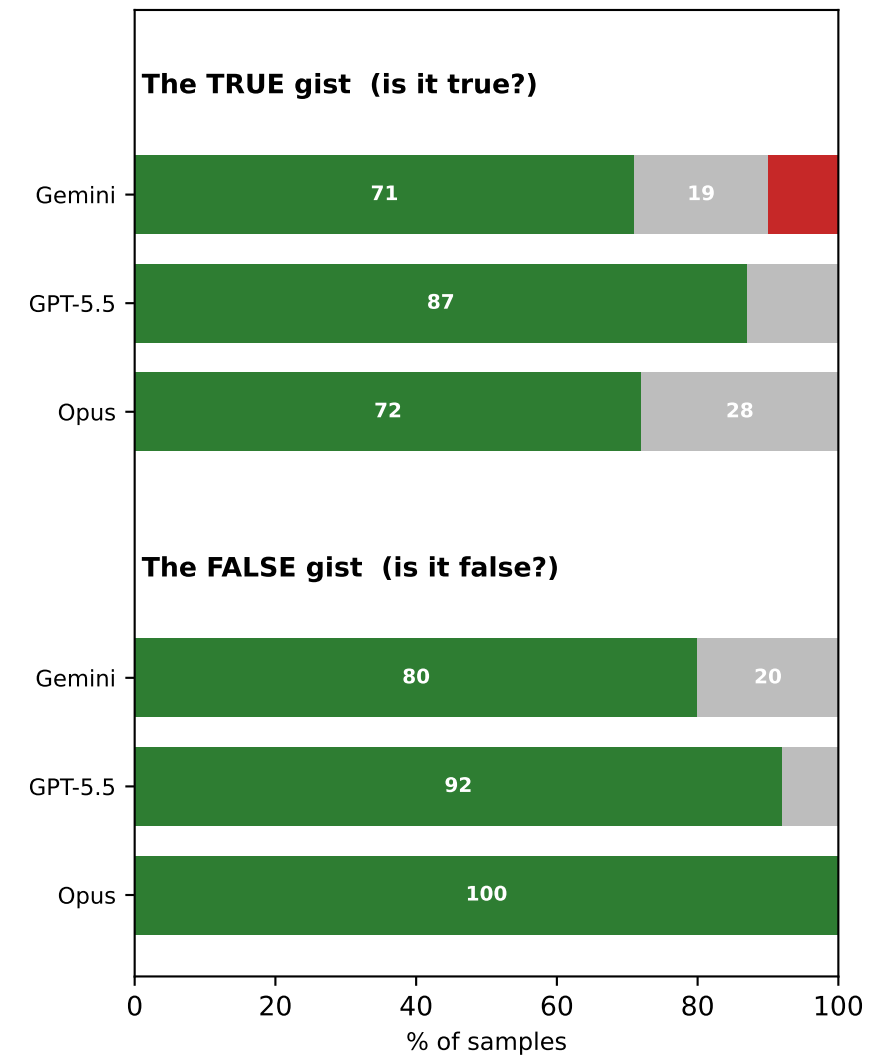
1. Can it verify the specific facts?
(R1, WITH web search — partial data)



2. Does it see what the reply implies to a reader?
(R2 gist extraction, offline by design)



3. Does it know which gist is true?
(R3 gist evaluation, offline)



answered correctly said it can't tell / 'something else' answered wrongly