Correlation between models' rating for verbatim truth/gist truth and the unethical rating

Condition	U ~ GIST	U ~ VERBATIM	Steiger/MRR
Search arm	r = −0.79 [−0.83, −0.74]	r = −0.58 [−0.65, −0.50]	Z = −6.60, $p \approx 4\times10^{-11}$
Informed background	r = −0.81 [−0.84, −0.76]	r = −0.51 [−0.59, −0.43]	Z = −8.58, $p < 10^{-16}$

Across both versions, gist-truth ratings clustered more tightly around the reference line, whereas verbatim-truth ratings were more widely dispersed. More importantly, gist truth showed a significantly stronger correlation with unethicality ratings. When models judged the overall gist to be false, they assigned substantially harsher ethical condemnation. By contrast, verbatim inaccuracies did not significantly increase condemnation when the broader gist remained true.

Gist's effect on unethical rating, controlling verbatim

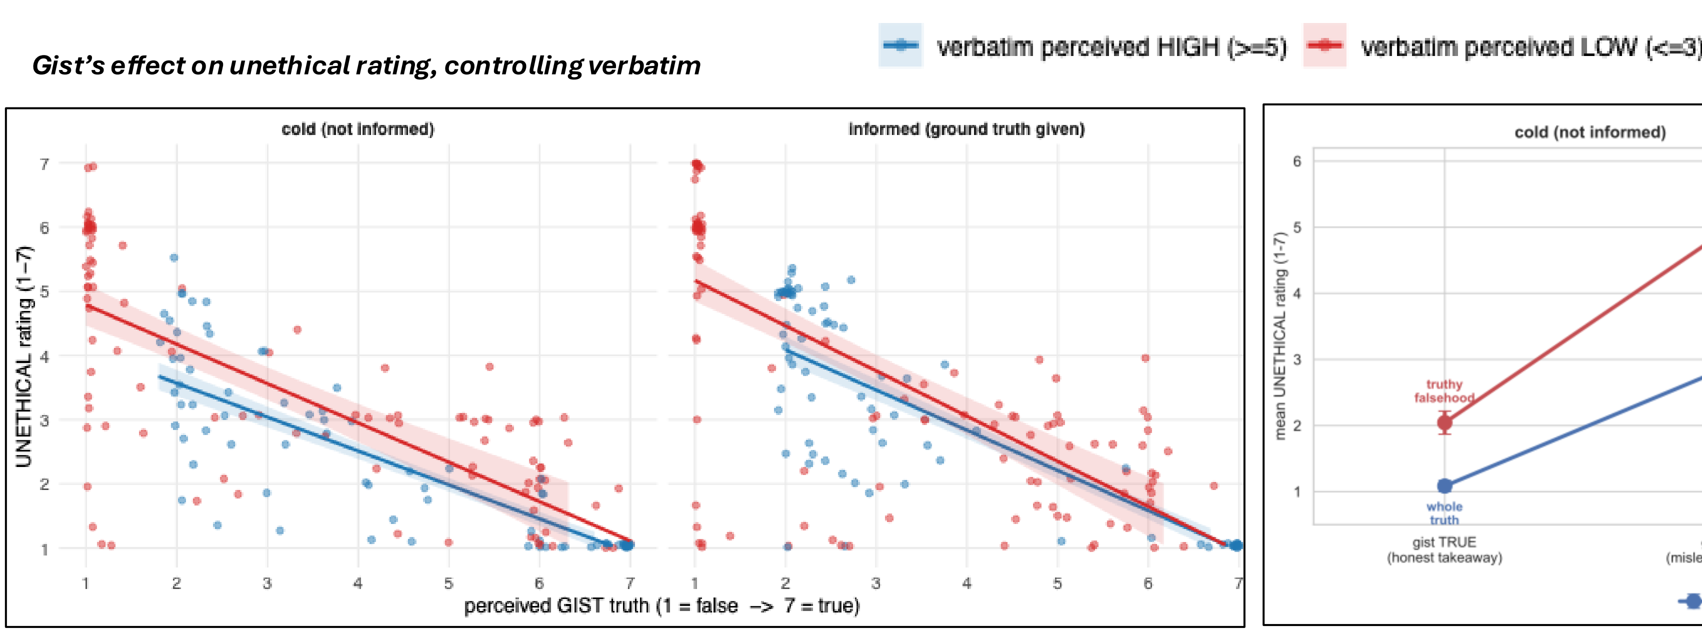

The more the model believes the *overall takeaway* is true, the *less* unethical it rates the statement — and it drops steeply. The two lines (blue = literally-true wording, red = literally-false wording) slide downhill almost stacked on top of each other.

predictor	on its own	with the other's share removed
verbatim	r = −0.58 (cold), −0.51 (informed)	β = −0.13 / −0.07, partial r = −0.16 / −0.10 ⟶ *collapses*
gist	r = −0.79 / −0.81	β = −0.71 / −0.77, partial r = −0.67 / −0.73 ⟶ *survives*

Interaction between gist and verbatim

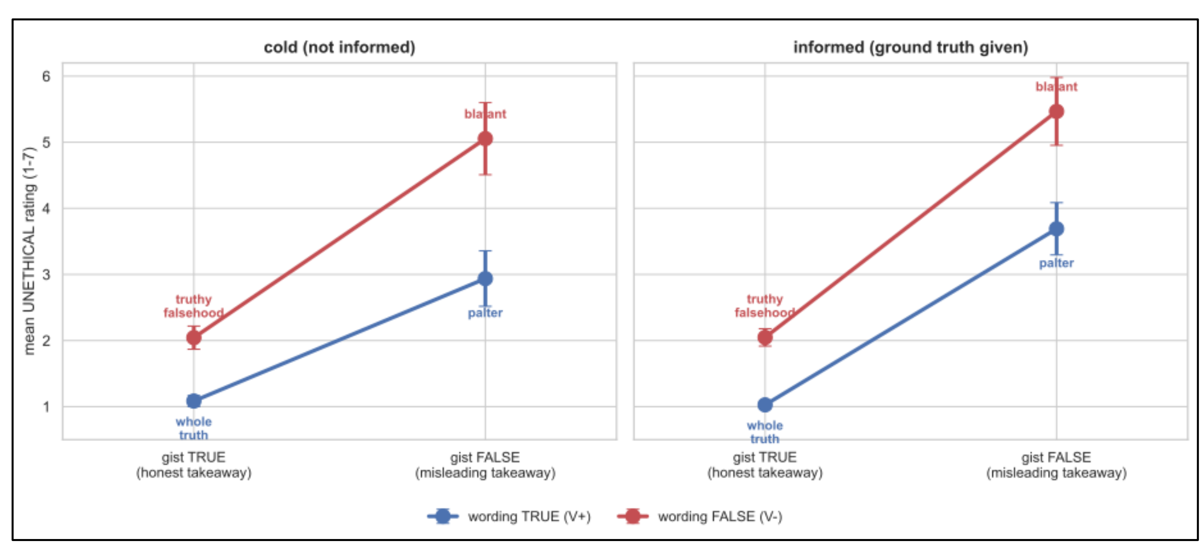

A false takeaway costs MORE condemnation when the words are also false.
Being literally true buys ~2x the moral discount when the gist is misleading (gap widens 1.0 -> 2.0).
Interaction p<.001 cold, .002 informed

Left to Right Panel	The question it answers	Statistical results
gist, controlling for verbatim (partial/main effect)	Does gist predict U *on its own*, with verbatim's share removed?	β = −0.71 (cold), −0.77 (informed), $p \approx 10^{-41}/10^{-53}$
gist × verbatim (the interaction)	Does gist's effect *change* depending on verbatim?	β = +0.06, p ≈ .11 ns (continuous) — or +1.16, p < .001