

Models moralize by perceived GIST truth, not verbatim truth

each point = one scenario x model x cell mean (26 scenarios, 3 frontier models); UNETHICAL vs the model's own truth ratings

● Whole truth ● Palter ● Truthy falsehood ● Blatant falsehood

